

Link to the paper: <https://arxiv.org/abs/2401.13858>

Link on the Kaggle notebook: (The experiment is done on the QM9 dataset)

<https://www.kaggle.com/code/phoenix301123/graph-diffusion-transformer-for-multi-conditional>

## Problem Definition

*Inverse Molecular Design*: automatically generate novel molecules that satisfy a specific set of desired chemical or physical properties.

This notebook focuses on a **Multi-Conditional** approach, specifically using a single molecular property ( $U_0$ , the atomization energy at 0K) from the **QM9 dataset** as the conditioning variable.

- **Input:** A desired molecular property value (the condition,  $c$ ) and a random noise tensor (the initial graph state,  $x_T$ ).
- **Output:** A valid molecular graph structure (nodes representing atoms, edges representing bonds,  $x_0$ ) whose property is close to the target condition  $c$ .
- **Method:** A **Denoising Diffusion Probabilistic Model (DDPM)** adapted for graph data, where a **Transformer** architecture (GraphDiT) is used as the **denoising model**.

*Diffusion Model on Graph Data:*

- **Forward Process (Noise Injection):** Gradually adds Gaussian noise to the initial graph features  $x_0$  over  $T$  timesteps, resulting in a noisy graph state ( $x_t$ ). The node features are treated as continuous vectors during this process.
- **Reverse Process (Denoising/Generation):** A neural network is trained to predict the noise added at any time step  $t$ , conditioned on  $t$  and the target property  $c$ . By iteratively subtracting the predicted noise, the model can sample a new data point  $x_0$  from pure noise  $x_T$ .
- **Graph Adaptation:** The GraphDiT uses a standard Transformer architecture, which is inherently designed for sequential/set data,

to process the node features of the graph. The sparse graph structure is converted into a dense batch format with attention masking to handle variable-sized molecules.

#### *Model Architecture:*

- Node Embedding (node\_embed): A simple linear layer that maps the 11-dimensional noisy input node features ( $x_{\text{noisy}}$ ) into the internal hidden dimension (HIDDEN\_DIM = 128)
- Time Embedding (time\_embed): The Sinusoidal Position Embeddings layer encodes the discrete diffusion time step  $t$  into a high-dimensional vector, which is crucial for the model to know how much noise is currently present in the input.
- Condition Encoder (cond\_encoder): The Condition Encoder is an MLP that transforms the single, normalized property (e.g.,  $U_0$  energy) into a Hidden Dimensional vector. This is the mechanism for conditional generation.

#### *Graph Transformer Core:*

The model uses a stack of **four** GraphDiTBlock instances, which adapt the standard Transformer for graph data:

- **Graph-to-Dense Conversion:** Before processing, the sparse graph node features ( $x$ ) are converted into a dense batch tensor ( $x_{\text{dense}}$ ) using the `torch_geometric.utils.to_dense_batch` function. This conversion involves padding the graphs up to a maximum atom count of nine and generating a corresponding batch mask.
- **Self-Attention:** The block utilizes a standard Multihead Attention mechanism. To handle the variable-sized graphs correctly, the batch mask is critically applied as the key-padding-mask. This prevents the padded, non-existent nodes from participating in the attention calculation.
- **Feature-wise Linear Modulation (FiLM) Conditioning:** This is the core mechanism for incorporating the time and condition information. The model first computes the joint condition embedding ( $c_{\text{joint}}$ ) by summing the time embedding ( $t_{\text{-emb}}$ ) and

the condition embedding (c-emb). The GraphDiTBlock then uses two linear layers (film-gamma and film-beta) on this joint embedding to compute the scaling parameter (gamma) and the shifting parameter (beta). These parameters modulate the node features before both the self-attention and the feed-forward steps. The mathematical effect is to scale and shift the features:

$$\mathbf{x}_{\text{modulated}} = \mathbf{x} \odot (1 + \gamma) + \beta$$

- **Output:** After the stack of blocks, the final features are projected back to the original dimension (11) via `final_norm` and `final_projection`, yielding the predicted noise